Batch cost at matched request concurrency

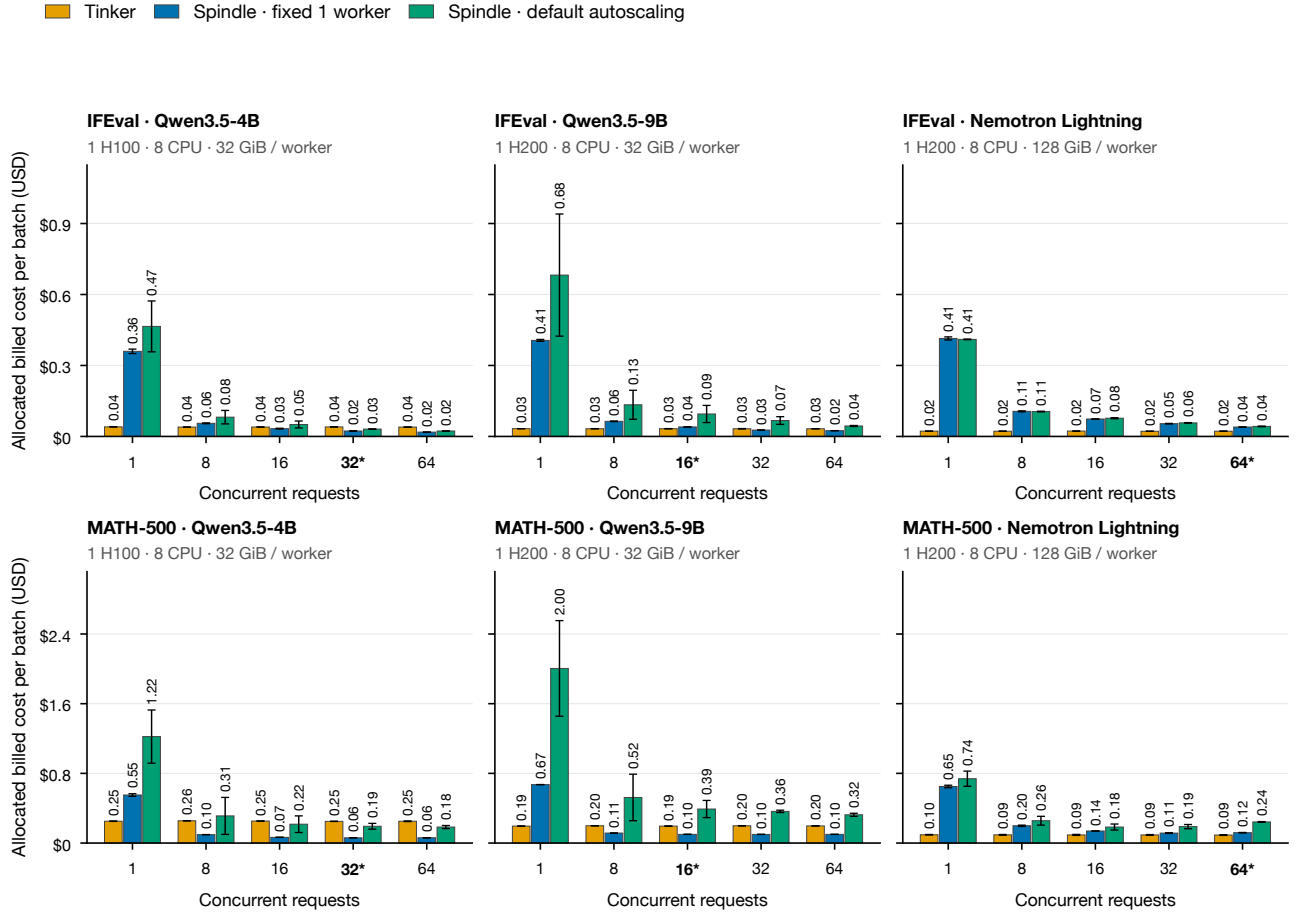

128 IFEval prompts; 64 MATH problems × 4 samples. Two-pass means; whiskers show range.

Billing: tokens / GPU time (auto: 1-min). Autoscaling includes GPUs retained across phases.

* Spindle default autoscaling target per worker; client concurrency was swept.