Warm throughput at matched request concurrency

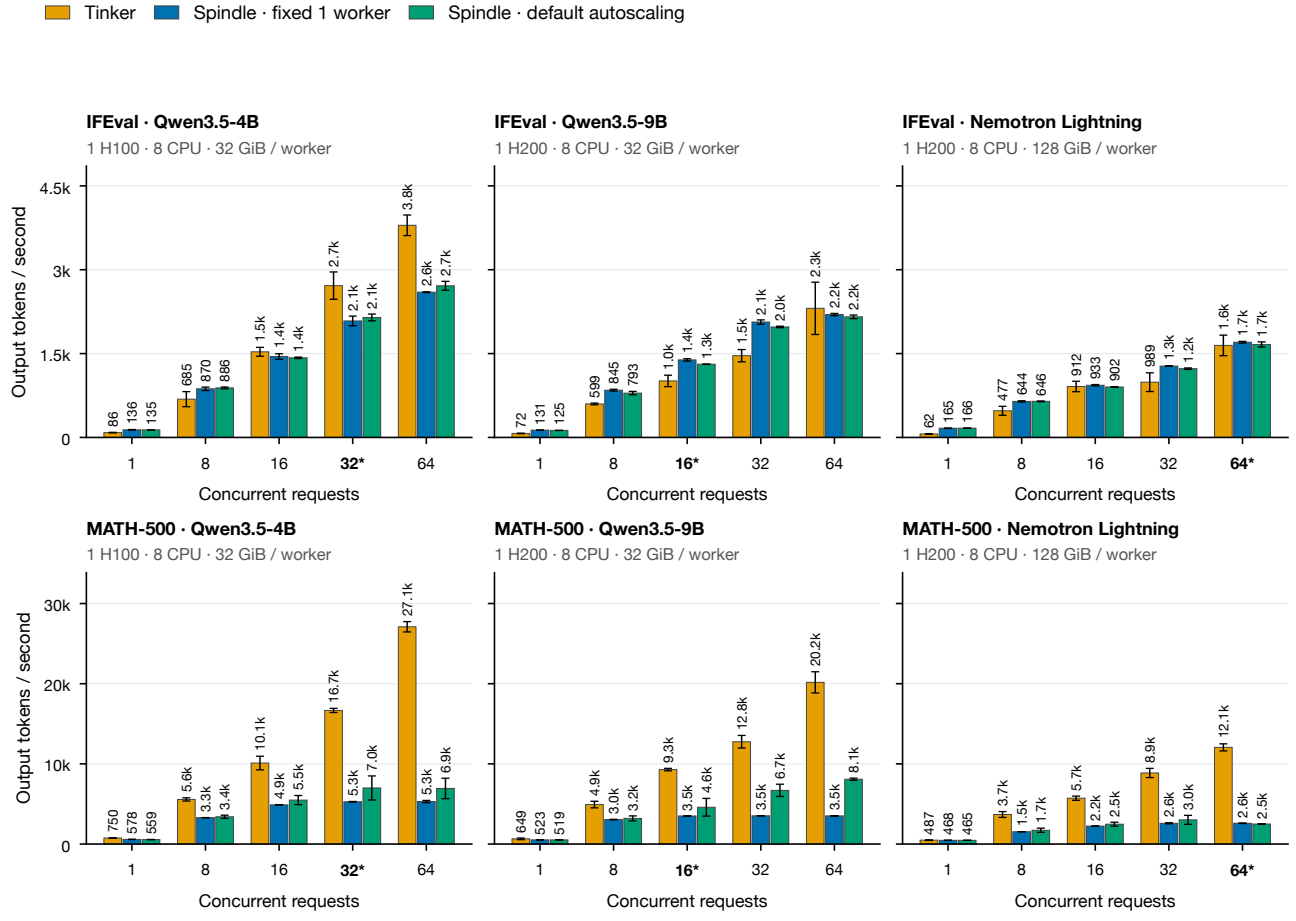

128 IFEval prompts; 64 MATH problems × 4 samples. Two-pass means; whiskers show range.

Warm API batches; scale-up during sampling included. Rows use different scales.

* Spindle default autoscaling target per worker; client concurrency was swept.